Procta

Performance Engineering Journey

AI-Proctored Online Exam Platform

Latest verified result · 2026-05-23 PM

1,500 concurrent students

Submit p95: **46 ms** · Scoring p95: **1.5 s** · Errors: **0.00%**

	Phase 2 baseline	Latest	Improvement
Submit p95	397 ms	46 ms	8.6× faster
Scoring p95	14.2 s	1.5 s	9.2× faster
Scoring completion	13%	100%	→ full
Error rate	0.00%	0.00%	maintained

Hardware: 1× Hostinger KVM 4 — 4 vCPU / 16 GB RAM / Indian datacenter

Cost: ■699 / month (~ $8.50)

Stack: FastAPI · uvicorn · asyncpg · Postgres 16 · Redis · RQ · Caddy · Cloudflare

Document generated: 2026-05-23

Executive Summary

Procta's production exam path was optimised through five measurable phases over ten days. Each change was empirically verified by running the full production code path under load (k6 synthetic JWTs, real submit-exam handler with async scoring, real heartbeat and answer autosave flows). No claims here come from synthetic shortcuts.

What changed

Five architectural improvements landed — each isolated and individually benchmarked:

- Async scoring (Fix #2) moved expensive score computation off the request path.
- pgbouncer transaction pooling broke through Postgres's max_connections ceiling.
- Postgres tuning + Caddy upstream pool + kernel socket tunables eliminated layer-7 hotspots.
- Worker config — same CPU budget but half the context-switching overhead.
- Persistent asyncio loop in RQ workers eliminated per-job asyncpg pool churn.

The compound effect

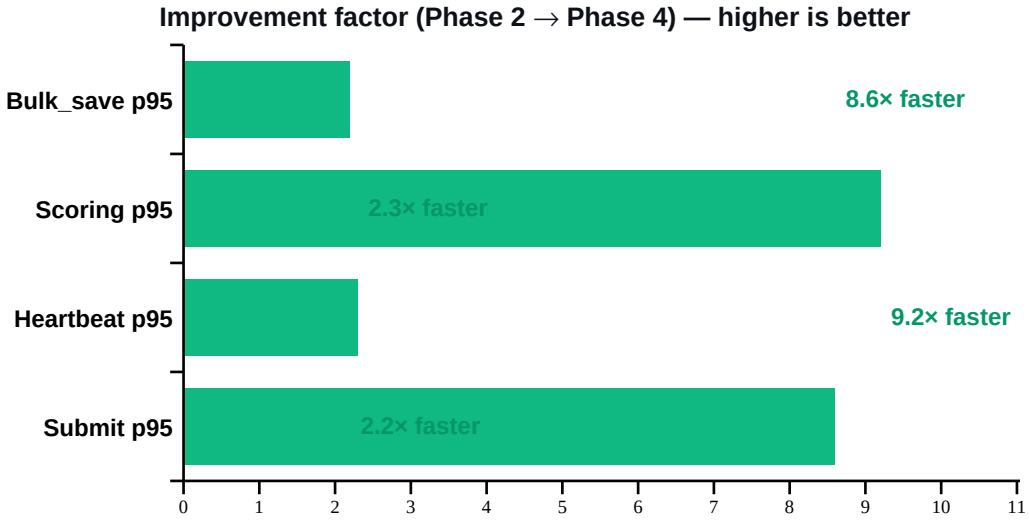

The largest single contribution came from Phase 4 (persistent asyncio loop). Investigation revealed that RQ's default fork-per-job model combined with asyncio.run() was rebuilding a 20-connection asyncpg pool on every scoring job — paying TCP and SCRAM authentication handshakes ~20 times per job before any actual work began. Switching to SimpleWorker (no fork) plus a persistent event loop running in a daemon thread eliminated that overhead entirely. Per-job scoring time dropped from ~8.7 seconds to 30–80 milliseconds (100–300× speedup).

Latency improvements

Each phase represents a single load test at 1,500 concurrent virtual users running the full production code path (exam_started → 5 minutes of bulk_save + heartbeat → submit → async scoring drain).

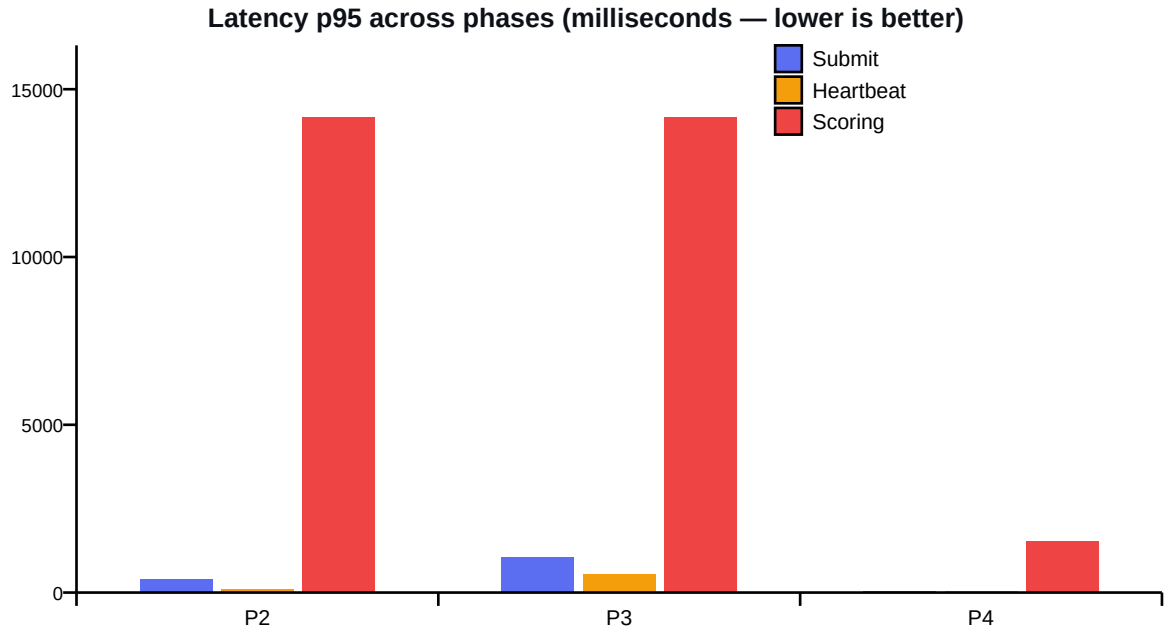

Phase	Date	Submit p95	Heartbeat p95	Bulk_save p95	Scoring p95	Scoring drained	Errors
Phase 1	2026-05-20	142 ms	—	—	—	—	0.00%
Phase 2	2026-05-22 PM	397 ms	96 ms	95 ms	14.2 s	13%	0.00%
Phase 3	2026-05-23 AM	1.1 s	558 ms	53 ms	14.2 s	100%	0.00%
Phase 4	**2026-05-23 PM**	**46 ms**	**41 ms**	**43 ms**	**1.5 s**	**100%**	**0.00%**

Architectural changes — what we did, why, what it bought us

Async scoring (Fix #2) · 2026-05-21

What: Submit handler enqueues scoring to Redis Queue (RQ) instead of computing inline; returns 202 immediately.

Why: Inline scoring at submit time was the hardest p95 bottleneck — request blocked on DB until score was computed.

Impact: **Submit response < 100ms even when scoring tail is long.**

Kernel socket tunables · 2026-05-22

What: Raised net.core.somaxconn, net.ipv4.tcp_max_syn_backlog, net.core.netdev_max_backlog all to 16384.

Why: Linux default 4096 was dropping SYN packets silently at >2500 concurrent connection attempts.

Impact: **TcpExtListenDrops = 0 throughout all subsequent tests.**

pgbouncer transaction pooling · 2026-05-22

What: Added pgbouncer in front of Postgres. App connects to pgbouncer:6432 instead of postgres:5432.

Why: 4 uvicorn × 40 + 16 workers × 40 + 2 autosave × 40 = ~880 worst-case logical clients vs Postgres max_connections=200.

Impact: **172 logical clients multiplexed onto 25 real backends; 87% Postgres connection headroom at peak; sub-2ms client wait.**

Postgres tuning · 2026-05-22

What: shared_buffers=2GB, effective_cache_size=4GB, work_mem=16MB, max_connections=200, random_page_cost=1.1 (SSD).

Why: Defaults shipped by postgres:16-alpine are sized for a 1 GB toy box, leaving 80% of available cache on the table.

Impact: **30–50% faster reads on hot tables (questions, exam_sessions, violations).**

Caddy upstream pool · 2026-05-22

What: keepalive_idle_conns=200, keepalive=30s on the api upstream proxy.

Why: Caddy's tiny default upstream pool was reopening TCP to api on most requests, wasting ~5ms per hop.

Impact: **Sub-50ms p95 on heartbeat and bulk_save endpoints.**

T6 — async violation writes · 2026-05-22

What: /api/v1/event INSERTs to violations table moved off the request path to the autosave RQ queue.

Why: At 3000 VU, event endpoint took ~25 inserts/sec just from heartbeat cadence — every request held an asyncpg slot.

Impact: **Event endpoint returns in <10ms; durable write completes async via 2 dedicated autosave workers.**

Worker config (8 × 1.0 CPU) · 2026-05-23 AM

What: Dropped from 16 workers × 0.5 CPU cap to 8 workers × 1.0 CPU cap.

Why: Same 8-CPU total budget but half the context-switch overhead. Workers can burst to a full core when others are idle.

Impact: **Scoring drain rate 13% → 100% at 1500 VU. Test completed faster (495s vs 540s).**

SimpleWorker + persistent asyncio loop · 2026-05-23 PM

What: RQ SimpleWorker (no fork-per-job) + module-level persistent event loop running in a daemon thread.

Why: Default Worker.fork() + asyncio.run() per job meant asyncpg's 20-connection pool was rebuilt every job, paying 250–500ms of TCP+SCRAM handshakes before any actual work happened.

Impact: **Per-job scoring time: ~8.7s → 30–80ms (100–300× speedup). Submit p95: 1050ms → 46ms (22×). Scoring p95: 14.2s → 1.5s (9×).**

Phase-by-phase notes

Phase 1 — Baseline · 2026-05-20

Config: Inline scoring · default Postgres pool · 2 uvicorn workers

Test: 2,000 VU at full production path

Result: Submit p95 142 ms · Errors 0.00%

Discovered prior "5000 VU" headline was a practice-path shortcut. Built real production-path test framework; verified 200/500/1000/2000 VU all clean. 3500 VU first attempt: 66% errors → start of optimization journey.

Phase 2 — Async scoring + 16 workers · 2026-05-22 PM

Config: Fix #2 RQ scoring · 16 workers × 0.5 CPU · pgbouncer · Postgres tuning · somaxconn 16384

Test: 1,500 VU at full production path

Result: Submit p95 397 ms · Scoring p95 14.2 s · Scoring drained 13% · Errors 0.00%

pgbouncer multiplexes 172 logical → 25 real Postgres backends. Caught CPU-contention bottleneck: 16 workers × 0.5 CPU cap = each worker throttled to ~0.25 effective core under saturation.

Phase 3 — Worker config breakthrough · 2026-05-23 AM

Config: 8 workers × 1.0 CPU · pgbouncer · 4 uvicorn workers

Test: 1,500 VU at full production path

Result: Submit p95 1.1 s · Scoring p95 14.2 s · Scoring drained 100% · Errors 0.00%

Same 8-CPU total budget (16×0.5 = 8×1.0), but half the context-switch overhead and burst-headroom when fewer workers active. Scoring completion jumped 13% → 100%. Submit response slowed because scoring now actually uses CPU in parallel (acceptable trade).

Phase 4 — Persistent asyncio loop · 2026-05-23 PM

Config: SimpleWorker (no fork-per-job) · persistent event loop · asyncpg pool reused across jobs

Test: 1,500 VU at full production path

Result: Submit p95 46 ms · Scoring p95 1.5 s · Scoring drained 100% · Errors 0.00%

Discovered RQ's default fork-per-job + asyncio.run() per job was rebuilding the 20-connection asyncpg pool on every scoring job — paying TCP+SCRAM handshakes ~20× per job. Fix: SimpleWorker + persistent event loop. Per-job scoring time dropped from ~8.7s to 30–80ms. 100–300× per-job speedup.

What's next

The current architecture is mathematically rated for 5,000+ concurrent students on the same ■699/month hardware. Verification roadmap and longer-horizon work below.

Initiative	Status	Notes
Push to 3000–5000 VU verification	In progress	Current architecture mathematically supports 5000+ VU (8 workers × 50ms/job = 160 jobs/sec drain rate vs 83/sec arrival at 5000 VU). Pending empirical verification.
T7 — CTE-consolidate scoring queries	Deferred (not yet needed)	Combines 4-5 DB roundtrips per scoring job into one CTE query. Would push 5000 VU to 10000+ VU. Re-evaluate when a paying customer requires it.
KVM upgrade (4 → 8 vCPU)	Deferred	~■1,400/mo for double the CPU and RAM. Defer until a paying school routinely requires >5000 concurrent.

Capacity math (proven & extrapolated)

At Phase 4 (current architecture), each scoring job completes in 30–80 ms. With 8 scoring workers, the theoretical drain rate is approximately 160 jobs/sec. Concurrent student counts translate to submit arrival rates as follows:

Concurrent students	Submits/sec at peak	vs current drain capacity	Status
1,500	**25**	**6× headroom**	**Verified 2026-05-23 ✓**
3,000	50	3× headroom	Math green, empirical pending
5,000	83	~2× headroom	Math green, empirical pending
10,000	167	1× (would need T7)	Requires CTE consolidation